

PM Caselet — LiveForever Care Brief

Bakul Badwal · GBUS 8636 Digital Product Management · Prof. Alexander Cowan · Individual assignment · v4, Sep 19, 2026 — revised after the Class 10 peer review (due Sun Sep 20, 11:59 PM) · Team of one on this subject

Live demo: bakulbadwal.github.io/liveforever-care-brief · Source: github.com/bakulbadwal/liveforever-care-brief (MIT) · Built on a private personal evidence engine I have run on my own data since 2021, and its OpenAI Build Week extension

THE ONE THING THIS BRIEF COMMITS TO

I do not yet know whether LiveForever is a personal evidence lab or a clinical handoff, and I know which observation settles it — and the one gate that has to open before it is worth observing. A person on a new medication can already see her own sleep and recovery data — the ring's app shows her everything. What she cannot do is put that data in front of the prescriber who has to decide, in a fifteen-minute follow-up, whether to keep, change or stop the dose. If the binding problem is the second one, the analysis engine is the wrong thing to keep improving. The test is whether she hands it over. The gate in front of it is whether any prescriber wants the page at all. Section 5 registers what dies on each outcome.


1 What's the basic idea for the product?

LiveForever Care Brief turns a person's own wearable and lab history into a one-page, source-linked before-and-after of one change, written for the clinician who has to decide what to do about it. Wearable recovery data (HRV, sleep, resting heart rate, readiness), a habit and medication log, lab history and cautious genomic context flow into a deterministic engine that answers *one* question — “since I started this medication, what happened to my sleep and recovery?” — as a 28-day before/after effect with its uncertainty, an evidence-quality grade, the missing evidence, and three or four questions for the prescriber. Python owns every number; the model explains, and never calculates, diagnoses or prescribes. The user reviews the draft and chooses whether to print it or export it. Nothing is transmitted; there is no account, no backend, no PHI in the public demo.

POSITIONING STATEMENT

TARGET For a person six weeks into a new prescription — a stimulant, an SSRI, a sleep aid — who wears a sleep and HRV ring and has a follow-up booked in the next thirty days, **NEED** who has to answer “how's it going?” from memory while a dose decision is made about her body, **CATEGORY** LiveForever Care Brief is a one-page pre-visit medication brief **BENEFIT** that shows the prescriber what the change actually did to her sleep and recovery — one effect, its uncertainty, an evidence grade, and three questions — in the minute the prescriber has.

ALTERNATIVE Unlike a screenshot of the ring app, or “I think I'm sleeping better,” **DIFFERENTIATION** it is written for the clinician, every number is source-linked and locked by code, and it says so when the evidence is too thin to act on.



The public demo on a fully synthetic persona and a lifestyle question (caffeine timing): the one question, evidence readiness (73 paired nights, 4/4 sources, 2 quality warnings, grade B), and the observed signal — +3.94 ms nightly HRV, 95% interval +1.13 to +6.42 — with the “values locked by code” boundary beside it. The wedge in this brief swaps the question for a medication change; the engine already runs 28-day before/after comparisons on logged interventions.

The gap — and the wedge. The data side of longevity is crowded and well funded: Oura, Whoop, Function Health, Superpower, InsideTracker, and every one of them makes the record bigger. None of them owns the moment the record has to survive: a scheduled follow-up with the person who wrote the prescription. **The wedge is the medication follow-up visit.** It has an unambiguous intervention (a drug started, titrated or stopped on a known date), a recipient with a decision at stake (keep, change, stop), a recurring cadence set by the prescriber rather than by curiosity (monthly during titration for stimulants, antidepressants and sleep medications, quarterly after), and outcomes — sleep, HRV, resting heart rate — that are exactly what a wearable measures well and exactly what those drugs move. I am this user: I bring my recovery score to a monthly medication follow-up with a physician assistant, from memory and a phone screen.

BUSINESS MODEL — PRICED PER VISIT, NOT PER MONTH, AND LEARNED BEFORE IT IS SET

The job recurs monthly during titration and one to four times a year after, so a subscription would price against how the job actually occurs. The candidate model is **one published price per brief**, the same for everyone, with the analysis itself free — the brief is the thing whose value is demonstrated at a visit, and the visit is the moment. **I am not setting the price in this document.** At single-digit briefs a week no price test clears the sample precondition; the price is learned from the twelve conversations in \$5 with one question — “*did you pay anyone to help you prepare for this visit, and what happened?*” — because the honest reference prices are a \$499 longevity-clinic membership on one side and a free screenshot on the other, and I do not yet know which one she compares me to. The model only exists under one reading of the product: if the reward is the handoff, a per-visit price is coherent; if the reward is understanding her own signal, this is a personal tool competing with free dashboards, and there is nothing to price per visit.

BUSINESS MODEL CANVAS — WHO PAYS, FOR WHAT, AND WHEN

Customer segments	Focal: patients in titration who wear a ring and have a follow-up booked. Declared, not focal: the prescriber (the veto) and the medical assistant (the gate to the chart) — the patient's orbit is drawn first; the recipient side is the first thing \$5 tests.
Value proposition	One page, one change, 28 days before against 28 after: effect, interval, grade, what is missing, three questions. Refusal when the evidence is thin — “there is nothing to prepare for” is a feature, not a failure.
Customer relationships	Self-serve, local-first, no account, nothing transmitted. The relationship the product serves is hers with her prescriber — a handoff, not a habit. The next visit is already booked; that is the only retention that exists.

Channels	The condition communities she is already in — r/ADHD, r/insomnia, r/ouraring — and longevity Substacks. Peer testimony after a visit (“the one-pager my PA actually read”): acquisition sits <i>downstream</i> of outcomes, storyboard panel 6.
Revenue streams	One published price per brief, analysis free, waived case by case — priced per visit because the job recurs per visit. Not set in this document: learned from the twelve conversations with one question, “did you pay anyone to help you prepare for this visit?” The honest reference prices are a free screenshot and a \$499 clinic membership.
How this answers the product/market-fit question	Every block is conditional on \$5. Under “handoff,” the price attaches to the visit and the channel is the community. Under “understanding,” it is a personal tool against free dashboards, with nothing to price per visit. The canvas is a bet the test decides, and fit is one observed behaviour: the page handed over.

2 Who's the target user, and what job are we hoping to do for them?

THE SCREENER — TWO FACTUAL QUESTIONS, NO MOTIVATION QUESTIONS

QUESTION	SCREENS IN
“Have you started, changed or stopped a prescription medication or a sleep supplement in the past 90 days — and were you wearing a sleep or HRV wearable most nights before the change?”	Yes to both
“Do you have a follow-up with the prescriber scheduled in the next 30 days?”	Yes

The first question fixes the intervention and guarantees a baseline; the second fixes the recipient and the date. Together they screen out the person who tracks out of curiosity — Cowan's question, *what's the intervention?*, has an answer before the conversation starts. I deliberately do *not* screen on which wearable, which drug, whether the prescriber is an MD, PA or NP, or how she feels about her care — a motivation question produces the answer you signalled you wanted, and both questions can be settled with a stranger in thirty seconds.

WHAT SHOES DOES SHE WEAR?

Running shoes, most days, with a ring on one hand. It is 9:40 on a weeknight; she is on the couch with the laptop, the ring app open on her phone, the patient portal open in another tab, and the follow-up reminder that arrived at lunch still unread in her inbox. Tomorrow has a 7 AM call. She started a new medication six weeks ago — a stimulant, an SSRI, a sleep aid, it does not matter which — and she has a follow-up in two weeks where the prescriber will ask “how's it going?” She has three years of ring data and knows her sleep changed, but she cannot say by how much, or whether the bad week was the drug or the travel. At the last visit she said “I think I'm sleeping better,” and the prescriber wrote that down.

WHERE DID SHE LAST READ THE NEWS?

On her phone: a longevity Substack, r/ouraring, and a condition-specific community — r/ADHD, r/insomnia — where people post their sleep scores after a dose change and ask each other what it means. Not a medical journal. That matters twice: the only channel that reaches her before a visit is a community she is already in, and the language she will believe is a peer's “here is what my prescriber actually looked at,” not a vendor's claim.

THE JOB

Show my prescriber what the medication actually did to my sleep and recovery — so the follow-up is about my data, not my recollection.

Trigger	The follow-up reminder, two weeks out — for a change that has a start date on the prescription label.
The clock	Days to two weeks — until the visit. After it, the brief has no recipient until the next one.
Frequency	Monthly during titration, then quarterly. A recurring appointment job, not a daily habit — and the one case in consumer health where “retention” means something: the next follow-up is already booked.
Where it sits	The evening before, at a laptop, with the ring app and the patient portal open in two other tabs.
Why it matters	A dose decision is being made about her body on the basis of “I think I'm sleeping better.” Years of data and no story is worse than no data: it produces the specific humiliation of knowing something is there and being unable to say it in a way the person with the prescription pad will act on.

ALTERNATIVES SHE IS USING TODAY — ALL REAL, AND EACH FAILS HER DIFFERENTLY

ALTERNATIVE	COST	HOW IT FAILS
Oura / Apple Health screenshots	Free	Shows her everything and gives the prescriber nothing to act on. This is the alternative that matters — it fully solves the job she <i>thinks</i> she has, which is “have the data with me.”
A ChatGPT summary of her exports	Free	Fluent, confident, and it will invent a number — the exact failure this product is built against. As Cowan put it in class, the DV of a general-purpose model is <i>the next word this person wants to hear</i> .
“I think I’m sleeping better”	Free	The modal choice, and what the chart currently records. Any proposition has to beat it, not just beat a screenshot.
Her own spreadsheet	Free	Stale by the visit, and it has no question in it — a prescriber cannot act on a sheet.
A longevity clinic (Function, Superpower, concierge)	\$\$\$	Replaces the prescriber instead of informing her; expensive; produces a bigger record, not a smaller handoff.

3 What's better *enough* about our proposition than the alternatives?

*If we hand a screened patient a one-page, source-linked before-and-after of her medication change — one question, one effect with its uncertainty, what is missing, and three questions for the prescriber — forty-eight hours before a scheduled follow-up, **then** she will hand it to the prescriber at the visit, unprompted — rather than reading it herself and answering “how's it going?” from memory as usual.*

The counterfactual is inside the hypothesis on purpose. “She will find it useful” is not usefully falsifiable; “*rather than reading it herself and answering from memory*” names the behaviour that means the reward was *understanding* — which is the other product, the personal evidence lab that already exists. As a user story: *As Maya, six weeks into a new medication with a follow-up in nine days, I want one page my prescriber can read in a minute that says what changed in my sleep and recovery, how sure I can be, and what to ask — so that I hand it over instead of describing it.*

WHEN AND WHY WE THINK IT'S BETTER

The delta over the alternatives is not analysis. Oura, Apple and a lab portal already give her every chart, free and better rendered than I will ever manage; the private engine already gives her the effect and the interval. If the proposition were “a better view of your own data,” I would be competing with the wearable makers on the job they already do, and I should stop. The claim is narrower and harder: the binding constraint is not *seeing*, it is getting a document into the hands of the person with the prescription pad, in a form that person will read in the time they have, before the visit is over. Better only inside two boundaries, stated rather than hedged: **(1) when a follow-up is scheduled** — outside that window the brief has no recipient and the honest output is “there is nothing to prepare for”; **(2) when the evidence clears readiness** — below the minimum paired sample or with the pre-change baseline missing, the correct output is a brief that says so. A prescriber's time spent on a grade-D signal is a cost I caused. Refusal is a feature here, not a failure.

WHY THE OBSERVABLE IS “HANDED OVER,” NOT “GENERATED”

The obvious DV is *briefs generated / records connected*. It cannot distinguish the two products this caselet is choosing between. If she generates a brief and keeps it, the reward was *understanding*, and the product is the evidence lab — for which the clinician questions, the print view and the export are dead weight. If she hands it over, the reward was *being taken seriously*, and the product is the handoff — for which the analysis depth is dead weight. A 70% generation rate would have read as “onboarding works” and retired the one uncertainty that decides what to build. **The verb has to be hers, and it has to happen in a room I am not in.**

TWO INTENTS COLLIDE — AND WHICH ONE WINS


Maya wants everything she noticed on the page. **The prescriber** — an internist, a psychiatric PA, a sleep-medicine NP; a queue, fifteen minutes, and a well-earned suspicion of patient-generated data — wants one question, one number with its uncertainty, and a reason to believe the number was not made up by a chatbot. They collide on length and on what counts as evidence. **The prescriber wins.** A brief that is not read is handed over and set aside, which scores as a handoff and produces nothing — so the recipient's constraints are a dependency of the experiment, not a later concern. That is also why §5 starts with the prescriber and not the patient. The price is real: I am overriding the user's instinct about her own history, so “brief abandoned before export” is a guardrail rather than a hope.

CERTAINTY PROPORTIONED TO THE DECISION — THE DESIGN PRINCIPLE BEHIND THE PAGE

The engine reports 95% intervals because that is the convention. The page does not treat $p = 0.05$ as the bar, because the decision at the visit is not one decision. “Keep taking it and re-check in four weeks” needs far less certainty than “stop a medication,” and a grade-B association is plenty for the first and never enough for the second. So the brief states the evidence grade *next to* the decision it is meant to inform, says which decisions it cannot inform, and leaves the dose to the prescriber. A fixed significance threshold is, in Cowan's phrase, a legal standard, not a clinical one — and the product's job is to let a regular person and a busy clinician talk their way through that proportioning without either having to know the statistics, while the numbers underneath stay rigorous and locked. That is also the market gap: general-purpose models optimise for the next word the user wants to hear; a tool that will tell you your signal is too weak to act on is a different product.

4 How does the customer experience unpack?

STORYBOARD — MAYA, 34, SIX WEEKS INTO A NEW MEDICATION, FOLLOW-UP IN FOURTEEN DAYS



CX MAP — DVS AND LINES IN THE SAND

STAGE	ACTIONABLE DV	CUTOFF	LINE IN THE SAND	OBSERVABLE TODAY?
Acquisition	records connected / demo opens, split by channel and by condition community	same session	> 30%	Yes — the only instrumented stage
Onboarding	briefs exported or printed / records connected, readiness reached (baseline present, grade ≥ C)	24 hours	> 60%	Partly — export is a local action; not reported
Engagement	handed to the prescriber at the visit / briefs exported	at the visit	> 50%	No — happens in a room the product is not in

Outcomes	dose decision made with the brief cited / briefs handed over (after-visit summary or portal note)	same visit	> 25%	No — in the prescriber's chart
Retention	a second brief before the next follow-up / first briefs	next visit (1–3 mo)	> 40%	No — but the next visit is already booked, so it is reachable

THE STAGE I AM LEAST SURE ABOUT — ENGAGEMENT, AND DELIBERATELY NOT OUTCOMES

Outcomes is the stage I can least *measure* — it lives in a chart I will never see — but I know what I want there, and the brief already refuses to claim it. Engagement is the stage where I do not know what the *reward* was, and the reward decides what the product is. If she exports the brief and it stays in her Downloads folder, the reward was understanding, and the product is the evidence lab that already exists; the prescriber questions and the print view are ornament. If she hands it over, the reward was being taken seriously, and the analysis depth is the ornament. The handoff row and this caselet's central question are the same question.

TWO THINGS THE MAP GETS WRONG ON PURPOSE

It is a loop, not a funnel. Acquisition sits downstream of Outcomes: the only scalable channel is a patient posting the page that worked, in the condition community listed as an alternative in §2. **And the recipient side is a second map I have not drawn.** Cowan's phrase for healthcare is an *orrery* of personas — patient, prescriber, clinic, payer — each with its own job and its own reward. This map draws the patient's orbit only. The prescriber's is the first thing §5 goes to find out, because a page nobody on that side wants is a dead product regardless of what the patient does.

5 What would I test next week?

WHICH KIND OF QUESTION — AND WHY IT DESERVES ATTENTION AHEAD OF THE OTHER TWO

It is a right-user-and-job question, stated plainly: *is the analysis the visible problem, or the binding one?* One question. A handoff has two ends, so the test has a gate on the recipient's side before it measures the patient's. It goes first because it is the question whose wrong answer wastes the most: every hour spent on the engine, the charts or the export path is wasted if the reward turns out to be the handoff, and every hour on the prescriber format is wasted if it turns out to be understanding. Ahead of **demand**, because a demand test presumes I know the job: a landing page for “make sense of your ring data” would convert — the quantified-self community converts on anything with a confidence interval in it — and would teach me nothing, because it tests the visible problem. Ahead of **usability**, because usability presumes the thing is worth using: watching someone connect a CSV tells me about the importer, not about whether a prescriber ever sees the result.

POSSIBLE VEHICLE	WHAT IT CAN TELL ME	WHY IT IS NOT THE ONE FOR THIS QUESTION
Interviews with patients	What she says she would do with a page; the vocabulary of the job	An interview cannot tell me whether anyone will act. Kept only as the post-visit diagnostic (“walk me through what happened with it”).
A landing page	How many click on “make sense of your ring data”	Counts interest in the <i>visible</i> problem. Cannot tell me whether I picked the right person or the right job.
A usability session on the demo	Where the importer and the export confuse people	Watching one person use it tells me about the interface, and presumes the page is worth handing over.
Prescriber conversations, then a concierge handoff	Whether the recipient exists; then whether the patient hands the page over, in a room I am not in	Chosen. The recipient conversations are the cheapest observation that can kill the thesis; the concierge handoff observes the behaviour itself, with the interface removed from the explanation. Together they fit the inference I need — a real decision about what to build — at the resources I have: one person, five weeks, \$0.

THE BIAS THIS TEST EXISTS TO CORRECT — I AM THE USER

LiveForever began as a private tool I built for myself and still use daily, and the wedge in this brief is literally my own monthly medication follow-up. That is the credibility — I know what it is like to answer “how’s it going?” from a phone screen — and it is the specific trap: an n-of-1 product built by its own user optimises for the visit *I* have, with the prescriber *I* have, who happens to be receptive. It has already shown up once: the first two versions of this product had no clinician in them at all. The recipient arrived late, in a hackathon extension, because a competition asked for a healthcare barrier. So this test starts on the other side of the room, with prescribers who are not mine — and then lets strangers’ prescribers, not mine, decide whether the page was worth forty seconds.

THE EXPERIMENT — ONE FOCAL TEST, WITH A RECIPIENT GATE IN FRONT OF IT

Gate — does the recipient exist? Week of Sep 14. Cost: five conversations, ~5 hours.	Five prescribers, none of them mine — starting with the two MD-MBAs in the product design class, then an NP and a PA through their referrals — each shown a real one-page brief built from my own data, and asked three questions: <i>Would you read this at a medication follow-up? For which patients — and which changes — would it change what you do? When should a patient send it, and when should they not?</i> Registered rule: ≥ 3 of 5 say they would read it for a medication follow-up → the test runs. ≤ 1 of 5 → the recipient does not exist at this level of generality; the handoff thesis dies here and the product is the evidence lab, without spending twelve patients to learn it. 2 of 5 → the answer is a <i>segment</i> , not a yes or no: run the test only on the patient type they named. This is the first-order answer Cowan said was cheap to get, and it is where the “which patients, when” question gets its first real data.
The test — does she hand it over? Sep 25 – Oct 20, as follow-ups fall.	Twelve screened patients with a follow-up in the next thirty days, briefs built by hand (Wizard of Oz — no new build). Recruited from r/ADHD, r/insomnia, r/ouraring and the Darden second-year class, which is unusually dense in people with a ring on one hand and a prescriber on the calendar. Each sends me an export — Oura or Apple Health — and the start date of the change. I build the 28-day before/after brief by hand in the engine, deliver it forty-eight hours before the visit as a single PDF, and say nothing about what to do with it . No “bring this,” no “your prescriber will love it.” What she does with it in a room I am not in is the measurement. A concierge test fits because it isolates the reward from the interface: nothing about the importer or the demo can explain the result, only whether a one-page, source-linked before-and-after is a thing a patient hands to a prescriber.
Focal DV	handed over / briefs delivered, at the visit — did she physically or digitally give it to the prescriber, unprompted, and can she evidence it (a photo of the annotated printout, a portal-message receipt, an after-visit summary that references it)? Observed behaviour, not stated intent. A brief she “meant to bring” is a miss — and if several land there, that is itself the finding.
Denominator	Ready briefs only — pre-change baseline present, minimum paired sample reached, grade $\geq$ C. Declared before the first conversation. Records that cannot reach readiness produce a brief that says so, and are reported under refusal correctness, not as failures.
Coverage, reported alongside	Screened-in patients attempted $\div$ screened-in. Never folded into the DV. Without it a high handoff rate is meaningless — I could have selected for people who already bring things to doctors.
Hard guardrail	Zero invented evidence, zero causal upgrades, zero dose language. One number in a brief that does not trace to the record, one “because” where the engine said “associated with,” or one sentence that reads as a dosing recommendation, voids that case and gets reported. A wrong number in a prescriber’s hands is worse than no brief.
Diagnostic, not decisional	After the visit, for every case: “ <i>walk me through what happened with it.</i> ” Not “was it useful.” Then one price question: “ <i>did you pay anyone to help you prepare for this visit, and what happened?</i> ”
Cost and time	Gate: ~5 hours, \$0. Test: ~44 hours over four weeks, \$0 cash — ~8 hours recruiting and screening, then ~3 hours per case (screen 30 min · ingest and clean the export ~1.5 hr · brief 30 min · post-visit conversation 30 min). Four weeks rather than two because the clock is the follow-up, not mine.

PASS / FAIL — REGISTERED BEFORE THE FIRST CONVERSATION

RESULT	READING	NEXT MOVE — AND WHAT GETS CUT
≥ 8 of 12 hand it over	The reward is being taken seriously . LiveForever is a clinical handoff with an evidence engine behind it.	Build: the handoff instrument (a share receipt); the medication before/after as the default brief; a prescriber-facing format tested with prescribers, not users; the portal-upload path; the source ledger as the product’s spine. Cut: the experiment planner, the interactive charts, PhenoAge on the front page, genomics anywhere a prescriber can see it.

≤ 3 of 12 hand it over	The reward is understanding . LiveForever is a personal evidence lab.	Build: more questions per record, the replication planner, the private daily loop. Cut: the Care Brief, the prescriber questions, the print and export path, the per-visit price, and the entire healthcare-handoff thesis — which is a different company, with a different buyer.
4 – 7	The brief is not the deciding variable . Neither product is validated.	A third named next move, not a shrug: re-run six, holding the visit type constant — a fifteen-minute primary-care follow-up and a thirty-minute psychiatric titration visit are not the same job, and the gate conversations will already have told me which one to hold.

WHY THE BANDS ARE THIS WIDE

A first draft would use $\leq 4 / \geq 7$, leaving only 5 and 6 undefined. A narrow band always returns a verdict, including from a coin flip. Binomial check on $n = 12$ against the two hypotheses this test is actually between — a 70% world and a 25% world — with a 50% world as the case the bands must *refuse* to call:

BANDS	P(FALSE PIVOT) TRUE 70%	P(FALSE PERSEVERE) TRUE 25%	P(FLAGS AMBIGUITY) TRUE 50%
$\leq 4 / \geq 7$	0.9%	1.4%	41.9%
$\leq 3 / \geq 8$ — adopted	0.2%	0.3%	73.3%

Under the narrow band a genuine 50/50 world is misread as a confident verdict 58.1% of the time. $n = 12$ is **enough** not because twelve is large but because the two hypotheses are far apart; the 4–7 band exists for the world where they are not. The same principle as the brief itself: the certainty is proportioned to the decision, which here is what to build next, not whether to launch a rocket.

WHY THIS IS A REAL DECISION

The two outcomes lead to **different code, a different buyer, a different roadmap, and a different business model**. Under “taken seriously,” the per-visit price is coherent, the prescriber’s format constraints become a dependency, and the next user research is with prescribers. Under “understanding,” there is nothing to price per visit and LiveForever is a personal tool in a market Oura and Apple own. And the gate can end it earlier still. I cannot write a next sprint that satisfies both readings.

THREE WAYS THIS TEST COULD LIE TO ME

- I am the concierge, and I know the data.** My hand-built brief will be better than the engine’s, and I will pick the question well. Mitigation: the patient names the change, not me; one PDF, no commentary; and I do not answer “should I bring this?” until after the visit.
- People who volunteer for a study about doctor visits already bring things to doctors.** That is the coverage number’s job, and why the screener is two factual questions with no motivation question in it.
- Handed over is not acted on.** A prescriber can accept the page and set it aside. Deliberately — I am measuring the handoff, not the decision; the decision lives in a chart I cannot read. That the handoff predicts action is carried as a stated belief, not a result, and I would rather have named that failure mode than discovered it.

ONE THING, NOT EVERYTHING — WHAT THIS BRIEF DELIBERATELY DOES NOT DISCUSS

More data sources · more wearables · EHR integration · a prescriber-side product · PhenoAge and biological-age marketing · the genomics layer · a subscription · a coaching or chat interface · anything called “AI health coach.” Every one of them lives to the right of the question in §5, and not one is resolved by anything I could do next week. Cowan’s instruction for this assignment was to set up a couple of yes-or-no questions and go through the process; this brief asks one — with a gate that can end it early — and says what dies on each answer.